


Linux Container The Hardway


Traditional Software Deployment - Dependency Hell



Traditional Software Deployment - Portability



Virtual Machine



- Advantages
 - Portable
 - Strictly Isolated
 - Multi-OS
- Disadvantages
 - Time consuming
 - Resource usage
 - Startup time
 - Scalability

What is container?



What is container?

An environment for executing processes with configurable isolation and resource limitations.


<https://github.com/opencontainers/runtime-spec/blob/main/glossary.md#container>

Why do we need container(s) ?

.... The goal of a Standard Container is to encapsulate a software component and all its dependencies in a format that any compliant runtime can run it without extra dependencies, regardless of the underlying machine and the contents of the container.

<https://github.com/opencontainers/runtime-spec/blob/main/principles.md#the-5-principles-of-standard-containers>

Linux Container & Virtual Machine



- Pros
 - Efficiency resource usage
 - Faster startup
 - Scalability
- Cons
 - Less secure than VM?
 - Same as Host OS

- Unix chroot system - 1979
- Free BSD Jails - 2000
- Linux VServer - 2001
- OpenVZ - 2005
 - ... Introduction of control groups (2006)
- Linux Container (LXC) - The most stable container system (2008)
- Cloud Foundry Warden - 2011


... The Golden Age

- Docker - 2013
- Windows support running (Windows & Linux) containers on Windows OS families

Why is Docker so popular ?

- Easy to use
- Good Documentations
- Tooling & CLI
 - Image Management (docker build/push/pull)
 - Public Image Registry
 - Standardize Image Build with Dockerfile
- Ecosystem and Cloud Support
 - Docker Swarm - An orchestration tool to manage clusters of Docker containers
 - Docker Trusted Registry - A private registry for trusted Docker images
 - Docker Compose - A tool for launching applications with numerous containers that need to exchange data.
 - Docker Desktop - A tool for creating Docker-enabled virtual machines.

Linux Container - Main Components



Namespaces

Namespaces

- Namespacing is a mechanism to provide processes with different view on different system resources
- It's a kind of process isolation
- Namespaces kind
 - Mount
 - Process ID
 - UTS
 - User
 - Network
 - CGroup
 - IPC (Inter-Process Communication)
 - Time
- Working with namespace
 - Create new namespace: **unshare**
 - Enter a namespace: **nsenter**

Demo Namespaces in Linux

Filesystem Isolation

*A **chroot** on Unix and Unix-like operating systems is an operation that changes the apparent root directory for the current running process and its children. A program that is run in such a modified environment cannot name (and therefore normally cannot access) files outside the designated directory tree*

<https://en.wikipedia.org/wiki/Chroot>

Control Group (CGroup)

Control Group

- `cpuset` - assigns individual processor(s) and memory nodes to task(s) in a group;
- `cpu` - uses the scheduler to provide cgroup tasks access to the processor resources;
- `cpuacct` - generates reports about processor usage by a group;
- `pid` - sets limit to number of processes in a group.
- `io` - sets limit to read/write from/to **block devices**;
- `memory` - sets limit on memory usage by a task(s) from a group;
- `devices` - allows access to devices by a task(s) from a group;
- `freezer` - allows to suspend/resume for a task(s) from a group;
- `net_cls` - allows to mark network packets from task(s) from a group;
- `net_prio` - provides a way to dynamically set the priority of network traffic per network interface for a group;
- `perf_event` - provides access to **perf events** to a group;
- `hugetlb` - activates support for **huge pages** for a group;

Interact with Control Group

Working with filesystem

```
root@basic:/home/vagrant# findmnt | grep cgroup
|-/sys/fs/cgroup
  |-/sys/fs/cgroup/unified
  |-/sys/fs/cgroup/systemd
  |-/sys/fs/cgroup/perf_event
  |-/sys/fs/cgroup/net_cls,net_prio
  |-/sys/fs/cgroup/cpu,cpuacct
  |-/sys/fs/cgroup/rdma
  |-/sys/fs/cgroup/devices
  |-/sys/fs/cgroup/memory
  |-/sys/fs/cgroup/hugetlb
  |-/sys/fs/cgroup/freezer
  |-/sys/fs/cgroup/pids
  |-/sys/fs/cgroup/blkio
  |-/sys/fs/cgroup/cpuset
```

Interact with Control Group

Helpers from *cgroup-tools*

```
root@basic:/home/vagrant# cgexec --help
Usage: cgexec [-h] [-g <controllers>:<path>] [--sticky] command [arguments] ...
Run the task in given control group(s)
-g <controllers>:<path>          Control group which should be added
-h, --help                         Display this help
--sticky                           cgred daemon does not change pidlist and children tasks
root@basic:/home/vagrant# cgcreate --help
Usage: cgcreate [-h] [-f mode] [-d mode] [-s mode] [-t <tuid>:<tgid>] [-a <agid>:<auid>] -g <controllers>:<path> [-g ...]
Create control group(s)
-a <tuid>:<tgid>                Owner of the group and all its files
-d, --dperm=mode                  Group directory permissions
-f, --fperm=mode                  Group file permissions
-g <controllers>:<path>          Control group which should be added
-h, --help                         Display this help
-s, --tperm=mode                  Tasks file permissions
-t <tuid>:<tgid>                Owner of the tasks file
```


Container Networking

Which problems do we need to solve?


- How to make containers think each of them has a dedicated network stack?
- How container communicate with each others & with host
- How to reach the outside world (e.g. the Internet) from inside the container?
- How to reach containers from the outside world (aka port publishing)?

Hand-On


Container Network - Dedicated Networking Stack




Container Network - Communication With Host




Container Network - Communication With Other Containers




Container Network - Communication With Other Containers




Container Network - External communication



Container Network - Port Forwarding



Container Network - Port Forwarding



CNI - Container Network Interface

Why do we need CNI ?

Container Runtime



CNCF Graduated



cri-o

CNCF Incubating




Firecracker



INCLAVARE



kata



lxd



rkt



OPEN CONTAINER INITIATIVE
runc



Singularity



SmartOS




Sysbox



WasmEdge runtime

Why do we need CNI ?

Cloud Native Network




CNI - Container Network Interface

The CNI project has 2 major parts

- The CNI specification documents
 - Go Libraries - <https://github.com/containernetworking/cni/tree/main/libcni>
 - A specification -
<https://github.com/containernetworking/cni/blob/main/SPEC.md>
- A set of based plugins - <https://github.com/containernetworking/plugins>
 - Interface management: create network interface/bridge/vlan
 - IP Address allocation: allocate IP address interfaces
 - Others chained plugin
 - tuning : tuning sysctl network parameters
 - portmap : iptables based port mapping
 - bandwidth : bandwidth-limiting for ingress/egress traffic
 - firewall: using iptables to allow/deny traffic to/from container
 - sbr: doing source-based routing


How CNI fit into the picture?



Demo - Using CNI with Docker

Union Mount

What the heck is it?



Yeah, I knew Union Mount ... but Why?

- **Contain images are quite big**
 - *It's expensive to allocate that much space every time we'd like to create a container from these images*
- **Thanks to union filesystem(s)**
 - *we only needs to create thin layer on top of the image and rest of it can be shared between all the containers.*
- **Reduce startup time**
 - *as there's no need to copy the image files and data*
- **Copy-on-write strategy if you want to modify those shared read-only files**

How does it implement in reality


- UnionFS
 - The original union filesystem
 - UnionFS doesn't seem to be actively developed anymore
- AUFS
 - A re-implementation of original UnionFS that added many new features....
 - But was rejected for merging into mainline Linux kernel. (*Its code was criticized for being "dense, unreadable, [and] uncommented"*)
 - Was default driver for Docker on Ubuntu/Debian but was replaced by OverlayFS
- OverlayFS
 - This is the filesystem used by default overlay2 Docker driver
 - It generally has better performance then aufs and has some nice features such as page cache sharing.
- ZFS
 - Created by Sun Microsystems (now Oracle)
 - Can not be shipped as part of Linux kernel because of licensing.
 - You could however use the ZFS on Linux (ZoL) project, but not ready for production.
- Btrfs
 - Btrfs is a default filesystem of Fedora 33.
 - It also has some useful features such as block-level operations, defragmentation, writeable snapshots and a lot more.

How Docker container filesystem looks like?

Toy Container Runtime

Container Images

Container Image



Container Image

- Layers
 - Container filesystem is split into multiple parts
 - Each layer is store in a single gzipped tarball
- Container Configuration: Contain metadata about images
 - OS/Architecture
 - Exposed ports/Entrypoint/Volumes/Env Variables/Labels/WorkingDir/User/Env Vars
 - A ordered list of digests (usually SHA-256) of **uncompressed** layer tarballs
 - A list of statement/command used to build the image - image history
- Blobs
 - Artifacts store in image registry (layers & manifest & configurations)
- Digests
 - Hash value of blobs
 - Used to uniquely identify blobs
 - Docker image use SHA-256 (OCI format support SHA-512 also)
 - Example: `sha256:09addbcf0db5a11803f29bddbdf31adce7e40d68750359f9a4eb4dcc54078f`
- Manifest: define the components that make up an image on a container registry
- Manifest list (fat manifest): which references image manifests for platform-specific versions of an image.

Image Manifest

Image Manifest V2, Schema 2

```
{  
  "schemaVersion": 2,  
  "mediaType": "application/vnd.docker.distribution.manifest.v2+json",  
  "config": {  
    "mediaType": "application/vnd.docker.container.image.v1+json",  
    "size": 8525,  
    "digest": "sha256:2b7ca628da40dc0cdcf6bf518b0f8cc0c5c3b83e89960e265e540cc7502ceee5"  
  },  
  "layers": [  
    {  
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",  
      "size": 54941777,  
      "digest": "sha256:6aefca2dc61dcbcd268b8a9861e552f9cdb69e57242faec64ac120d2355a9cla"  
    },  
    {  
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",  
      "size": 5155716,  
      "digest": "sha256:967757d5652770cfa81b6cc7577d65e06d336173da116d1fb5b2d349d5d44127"  
    },  
    {  
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",  
      "size": 10874928,  
      "digest": "sha256:c357e2c68cb3bf1e98dcb3eb6ceb16837253db71535921d6993c594588bffe04"  
    },  
    {  
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",  
      "size": 54578663,  
      "digest": "sha256:c766e27afb21eddf9ab3e4349700ebe697c32a4c6ada6af4f08282277a291a28"  
    },  
  ]  
}
```

OCI Image Manifest

```
{  
  "schemaVersion": 2,  
  "mediaType": "application/vnd.oci.image.manifest.v1+json",  
  "config": {  
    "mediaType": "application/vnd.oci.image.config.v1+json",  
    "size": 7023,  
    "digest": "sha256:b5b2b2c507a0944348e0303114d8d93aaaa081732b86451d9bc1f432a537bc7"  
  },  
  "layers": [  
    {  
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",  
      "size": 32654,  
      "digest": "sha256:9834876dcfb05cb167a5c24953eba58c4ac89b1adf57f28f2f9d09af107ee8f0"  
    },  
    {  
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",  
      "size": 16724,  
      "digest": "sha256:3c3a4604a545cdc127456d94e421cd355bca5b528f4a9c1905b15da2eb4a4c6b"  
    },  
    {  
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",  
      "size": 73109,  
      "digest": "sha256:ec4b8955958665577945c89419d1af06b5f7636b4ac3da7f12184802ad867736"  
    },  
  ],  
  "annotations": {  
    "com.example.key1": "value1",  
    "com.example.key2": "value2"  
  }  
}
```

Image Manifest List

Docker Image Manifest List v2

```
{  
  "manifests": [  
    {  
      "digest": "sha256:76f9e442fdf5c12efb5949407b0bb7ad28a413b8a5387a4243b1d43a14654060",  
      "mediaType": "application/vnd.docker.distribution.manifest.v2+json",  
      "platform": {  
        "architecture": "amd64",  
        "os": "linux"  
      },  
      "size": 2218  
    },  
    {  
      "digest": "sha256:04c13ef2cc265f31aa90decf9ba111f715b82d8c80dc2f84b13b1bce31b22f77",  
      "mediaType": "application/vnd.docker.distribution.manifest.v2+json",  
      "platform": {  
        "architecture": "arm64",  
        "os": "linux",  
        "variant": "v8"  
      },  
      "size": 2218  
    }  
  ],  
  "mediaType": "application/vnd.docker.distribution.manifest.list.v2+json",  
  "schemaVersion": 2  
}
```

OCI Image Index

```
{  
  "schemaVersion": 2,  
  "mediaType": "application/vnd.oci.image.index.v1+json",  
  "manifests": [  
    {  
      "mediaType": "application/vnd.oci.image.manifest.v1+json",  
      "size": 7143,  
      "digest": "sha256:e692418e4cbaf90ca69d05a66403747baa33ee08806650b51fab815ad7fc331f",  
      "platform": {  
        "architecture": "ppc64le",  
        "os": "linux"  
      }  
    },  
    {  
      "mediaType": "application/vnd.oci.image.manifest.v1+json",  
      "size": 7682,  
      "digest": "sha256:5b0bcabd1ed22e9fb1310cf6c2dec7cdef19f0ad69efa1f392e94a4333501270",  
      "platform": {  
        "architecture": "amd64",  
        "os": "linux"  
      }  
    }  
  ],  
  "annotations": {  
    "com.example.key1": "value1",  
    "com.example.key2": "value2"  
  }  
}
```

Meeting Container Registry API V2

Dockerfile Good Practices

Dockerfile Good Practices

- Using multi-stages Dockerfile
- Using '.dockerignore'
- Using Dockerfile linter
 - hadolint: <https://github.com/hadolint/hadolint>
- Using COPY/ADD options to change file metadata (permissions/owner)
- Write Cache Friendly Dockerfile
- Using BuildKit cache mount
- Using secrets/ssh mount